

Aggregation Semantics for Prioritization in Cooperative Automated Decision Making

Pavel Novoa-Hernández^{1,*} , Mariia Godz²  and David A. Pelta² 

¹ Dept. Computer Engineering and Systems, Universidad de La Laguna; pnovahe@ull.edu.es

² Dept. Computer Science and Artificial Intelligence, Universidad de Granada; {mariagodz, dpelta}@ugr.es

* Correspondence: pnovahe@ull.edu.es

Abstract

To be rewritten.

Keywords: cooperative automated decision making; aggregation operators; multi-criteria decision making; context-aware decision support; prioritization

1. Introduction

1.1. Research Questions and Contributions

2. Background and Related Work

2.1. Context-Aware and Cooperative Decision Support Systems

The increasing adoption of Artificial Intelligence (AI) in decision-support systems has highlighted the need to complement predictive performance with contextual, organizational, and domain-specific knowledge. In many real-world settings, decisions cannot be based solely on statistical predictions because institutional policies, expert knowledge, regulatory requirements, and operational constraints must also be considered. As a result, a growing body of research has focused on context-aware and cooperative decision-support systems capable of integrating heterogeneous sources of information into the decision-making process.

One important research direction concerns the incorporation of explicit knowledge representations alongside data-driven models. Early examples include cognitive-agent architectures that combine symbolic rules with machine-learning mechanisms to support context-aware decision making [1]. Similarly, ontology-based clinical decision-support systems have employed contextual ontologies and clinical pathways to select context-appropriate recommendations for individual cases [2]. More recent developments have extended these ideas through semantic models, knowledge graphs, and hybrid AI architectures that combine predictive models with structured domain knowledge to improve decision quality and interpretability [3,4].

A second research stream emphasizes cooperation between human and artificial intelligence. Rather than viewing decision support as a purely predictive task, these approaches consider decision making as a collaborative process in which machine-learning models provide evidence while human experts contribute contextual understanding, oversight, and governance. Cronholm and Göbel [5], for example, argue that effective decision support requires appropriate mechanisms for combining human and artificial intelligence rather than simply improving predictive accuracy. Likewise, recent surveys on human-centred AI highlight transparency, accountability, contextual awareness, and governance as fundamental requirements for trustworthy decision-support systems [6].

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *Mathematics* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](#) license.

Related ideas also appear in cooperative and distributed decision-making environments. Multi-agent systems, collective-intelligence frameworks, and cooperative AI architectures frequently rely on multiple decision sources that must be reconciled before a final recommendation is produced [7]. In these settings, contextual information often plays a coordinating role by constraining, guiding, or validating the outputs generated by predictive components. Similar concerns arise in governance-oriented decision-support systems, where contextual information may encode operational procedures, organizational priorities, safety requirements, or regulatory obligations [4,8].

Taken together, these studies clearly demonstrate the importance of contextual knowledge in decision-support systems. However, their primary focus typically lies in knowledge representation, ontology design, human–AI interaction, governance mechanisms, or system architecture. Once predictive evidence and contextual assessments have been generated, the numerical integration process that combines both sources of information is often treated as a secondary implementation detail. Consequently, relatively little attention has been devoted to understanding how alternative integration mechanisms affect the final recommendations produced by the system.

This limitation becomes particularly visible in architectures that explicitly separate predictive evidence from contextual assessment. Recent work on Cooperative Automated Decision-Making Systems (CADEMAS) [9] and its machine-learning instantiation CADEMAs-ML [10] exemplifies this design philosophy. CADEMAs organizes heterogeneous automated decision-making components and contextual knowledge into distinct layers, allowing predictive evidence and contextual assessments to be generated independently before being combined into a final recommendation. Such architectures make the integration mechanism itself an explicit design choice rather than an implicit implementation detail. As a result, the aggregation operator determines how predictive evidence interacts with contextual requirements, organizational priorities, or policy constraints, thereby shaping the semantics of the resulting decision process.

This observation suggests that aggregation should not be viewed merely as a numerical post-processing step. Instead, it can be interpreted as a governance mechanism that regulates the interaction between predictive and contextual information. Understanding the implications of different aggregation semantics therefore becomes essential for the design of cooperative and context-aware decision-support systems. The next subsection reviews the aggregation literature from this perspective.

2.2. Aggregation Operators and Decision Semantics

Aggregation operators constitute one of the central building blocks of multi-criteria decision making (MCDM), fuzzy decision support, and collective decision-making systems. Their purpose is to combine multiple sources of information into a single assessment while preserving, to a greater or lesser extent, the decision semantics associated with the original criteria. Consequently, aggregation has been studied extensively from both theoretical and applied perspectives.

Among the most influential aggregation frameworks are Ordered Weighted Averaging (OWA) operators [11,12], t-norm and t-conorm based operators, geometric aggregation functions, and non-additive integrals such as the Choquet and Sugeno integrals. Collectively, these approaches provide a rich mathematical vocabulary for representing different attitudes toward compensation, conjunction, disjunction, and interaction among criteria. From this perspective, the choice of aggregation operator is not merely a numerical convenience but a modeling decision that determines how favorable and unfavorable evidence may compensate for one another.

A substantial body of literature has explored these properties in traditional decision-making settings. Reviews by Mardani et al. [13] and Csiszár [14] highlight the widespread use of OWA-based aggregation and related fuzzy operators across diverse application domains. More recent studies have continued to investigate how aggregation semantics influence consensus formation, robustness, criterion interaction, and ranking behavior. Examples include consensus-oriented aggregation in group decision making [15,16], robust Choquet-based decision models [17,18], and applications employing geometric and OWA-type operators to capture alternative decision attitudes [19,20].

Despite this methodological richness, most aggregation studies are formulated within classical MCDM settings, where the aggregated inputs represent multiple criteria, attributes, preferences, or expert judgments associated with the same decision problem. The primary objective is typically to identify the most appropriate operator for combining these criteria or to develop new aggregation mechanisms with desirable mathematical properties. Comparatively less attention has been paid to architectures in which the aggregated inputs originate from fundamentally different decision layers.

This distinction is particularly important in cooperative automated decision-making systems. In architectures such as CADEMAs-ML, the aggregation process does not combine multiple criteria describing an alternative. Instead, it reconciles two conceptually distinct sources of evidence: a predictive assessment generated by machine-learning models and a contextual assessment encoding organizational priorities, regulatory requirements, expert knowledge, or policy constraints. Under this interpretation, the aggregation operator effectively determines the extent to which predictive evidence may compensate for weak contextual support, or conversely, whether contextual constraints can restrict otherwise favorable predictions. The operator therefore acts as an implicit governance mechanism rather than merely a criterion-combination tool.

For this reason, the present work focuses on three representative aggregation regimes exhibiting fundamentally different compensation semantics. The weighted arithmetic mean represents a fully compensatory integration strategy in which predictive and contextual information may offset one another. The weighted geometric mean represents a partially compensatory strategy that penalizes strong imbalances between both sources of evidence. Finally, the minimum operator represents a non-compensatory strategy in which the weakest component determines the final assessment. Together, these operators span a spectrum ranging from highly permissive to highly restrictive integration behavior while remaining interpretable and widely understood within the aggregation literature.

From the perspective of this study, the objective is therefore not to propose a new aggregation operator nor to extend aggregation theory. Instead, aggregation operators are treated as alternative decision semantics within a cooperative automated decision-making architecture. The key question is how different compensation regimes influence policy compliance, prioritization behavior, robustness to predictive perturbations, and ranking stability when predictive and contextual information must jointly determine the final recommendation.

2.3. Research Gap and Positioning

The reviewed literature reveals two largely independent research streams. On the one hand, context-aware and cooperative decision-support systems have developed increasingly sophisticated mechanisms for incorporating contextual knowledge into predictive processes through ontologies, symbolic rules, knowledge graphs, human oversight, and governance-oriented architectures [1,2,5,6]. These studies demonstrate the importance of contextual information but typically focus on knowledge representation, rule selection, system architecture, or human–AI interaction. Once predictive and contextual assessments

have been generated, the mechanism responsible for their numerical integration is often treated as an implementation detail.

On the other hand, aggregation theory and fuzzy MCDM have developed a rich collection of operators with diverse compensation semantics, including OWA operators, geometric means, t-norms, Choquet integrals, and Sugeno integrals [11–14]. This literature provides extensive theoretical knowledge regarding compensation, criterion interaction, robustness, consensus formation, and preference aggregation. However, aggregation operators are generally studied as mathematical constructs or as components of multicriteria decision methods, where the aggregated inputs represent criteria, preferences, or expert judgments associated with the same decision problem.

Consequently, a gap exists between context-aware decision-support frameworks and aggregation theory. The former provides mechanisms for representing contextual knowledge but rarely investigates how different aggregation semantics affect the preservation of contextual constraints after integration. The latter offers powerful aggregation mechanisms but seldom studies their implications when heterogeneous decision signals, such as predictive machine-learning outputs and contextual assessments, must jointly determine operational recommendations.

From this perspective, the gap addressed in this work is not the absence of aggregation operators, nor the lack of mathematical knowledge about their properties. Rather, it is the lack of a systematic understanding of aggregation as a governance mechanism within cooperative automated decision-making architectures. In particular, relatively little attention has been devoted to questions such as whether contextual assessments can effectively act as veto mechanisms, how different compensation regimes influence policy compliance, how sensitive prioritization outcomes are to predictive overconfidence, or how ranking stability is affected by uncertainty in contextual information.

These observations motivate the following research questions:

- *RQ1. How do different aggregation semantics affect the interaction between predictive evidence and contextual assessments in a cooperative automated decision-making architecture?*
- *RQ2. Under what conditions can aggregation operators preserve or violate contextual constraints and veto-like policies?*
- *RQ3. How do alternative aggregation regimes influence prioritization robustness under predictive overconfidence and contextual uncertainty?*
- *RQ4. To what extent do different aggregation semantics alter the stability of the resulting rankings and intervention priorities?*

Unlike traditional fuzzy MCDM studies, the objective of the present work is not to propose a new aggregation operator, derive new algebraic results, or solve a classical multicriteria decision problem. Likewise, unlike many context-aware DSS frameworks, the focus is not on developing new mechanisms for representing contextual knowledge. Instead, aggregation operators are interpreted as alternative policy semantics within the CADEMAs architecture. By keeping both the predictive subsystem and the contextual subsystem fixed, the study isolates the aggregation mechanism as an independent design variable and examines its implications for policy compliance, robustness, and prioritization behavior. This positioning distinguishes the contribution from both aggregation theory and context-aware decision-support research.

2.4. Chronological Comparison of Related Work

Table 1 summarizes representative contributions discussed in the previous subsections and positions the present work with respect to the existing literature. The comparison focuses on four dimensions that emerge from the identified research gap: (i) the explicit integration of predictive and contextual information, (ii) the treatment of contextual con-

Table 1. Chronological comparison of representative related work. “Partial” indicates that a contribution addresses the dimension only indirectly (e.g., through rules, consensus thresholds, or architectural design) without a dedicated analysis of predictive–contextual score fusion.

Reference	Year	Main focus	Predictive + contextual integration	Policy / veto analysis	Aggregation robustness	Cooperative ADM
Teng and Tan [1]	2008	Context-aware cognitive agents	✓	Partial	No	Partial
Yao and Kumar [2]	2012	Ontology-based clinical DSS	✓	Partial	No	No
Merkert et al. [21]	2015	ML in DSS survey	Partial	No	No	No
Moral et al. [15]	2017	Consensus aggregation	No	No	Partial	No
Mardani et al. [13]	2018	Fuzzy MCDM review	No	No	No	No
Csiszár [14]	2021	OWA theory	No	Partial	Partial	No
Cronholm and Göbel [5]	2022	Human-centred AI DSS	✓	Partial	No	Partial
Yang et al. [17]	2022	Robust Choquet-based MCDM	No	Partial	✓	No
Liao et al. [18]	2023	Contextual Choquet models	Partial	No	✓	No
Trillo et al. [16]	2023	Consensus-based GDM	No	Partial	No	No
Fanfani et al. [6]	2025	Human-centred AI survey	✓	Partial	No	Partial
Present work	2026	CADEMAS aggregation semantics	✓	✓	✓	✓

straints or veto-like policies, (iii) the analysis of aggregation robustness, and (iv) the study of aggregation within cooperative automated decision-making architectures.

The table illustrates that prior research has typically focused either on context-aware decision-support architectures or on aggregation methods in MCDM and fuzzy decision making. Comparatively little attention has been devoted to aggregation as a governance mechanism that simultaneously integrates predictive evidence, contextual knowledge, policy constraints, and robustness considerations within a cooperative ADM framework.

3. CADEMÁS and CADEMÁS-ML

The Cooperative Automated Decision-Making System (CADEMÁS) framework was introduced by Novoa-Hernández et al. [9] as a general architecture for context-aware decision support in which multiple heterogeneous Automated Decision-Making (ADM) units cooperate to produce a unified recommendation, score, or decision. Rather than prescribing a specific decision model, CADEMÁS provides a modular structure that separates the generation of decision evidence from its contextual evaluation and final integration.

This separation is particularly relevant for the problem studied in the present work. Many decision-support systems combine predictive evidence with contextual information, but these elements are often intertwined within the same model or decision process. In contrast, CADEMÁS explicitly distinguishes three conceptually different functions: (i)

evidence generation, (ii) contextual assessment, and (iii) decision integration. Consequently, the aggregation mechanism can be analysed independently of both the predictive subsystem and the contextual knowledge representation.

Formally, a CADEMAs instance is represented by the tuple

$$\mathcal{S} = \langle I, U, O, K, T \rangle, \quad (1)$$

where I denotes the Input Manager, U the set of ADM units, O the Decision Integration Layer, K the Contextual Modulator, and T the universe of admissible data representations.

Given an input case, the ADM units independently generate decision evidence from the available information. In parallel, the Contextual Modulator evaluates the extent to which the same case satisfies organizational priorities, expert knowledge, regulatory requirements, or other contextual constraints. The outputs of both subsystems are then combined by the Decision Integration Layer to produce the final decision score.

An important specialization of this framework is CADEMAs-ML [10], in which the ADM units correspond to machine-learning models. In this setting, multiple predictive models cooperate to produce a consolidated predictive assessment, while contextual information is represented through a fuzzy inference system inspired by Pérez-Cañedo et al. [22] and Pérez-Cañedo et al. [23]. The contextual subsystem translates organizational priorities and institutional requirements into a normalized degree of contextual satisfaction.

To focus on the role of aggregation, the internal implementation of both subsystems is abstracted in terms of two normalized scores. For each decision case i , the predictive subsystem produces a predictive score

$$R_i \in [0, 1], \quad (2)$$

which summarizes the cooperative evidence generated by the predictive ADM units. Depending on the application domain, R_i may represent a risk estimate, a confidence score, a probability of occurrence, or any other normalized predictive assessment.

Independently, the contextual subsystem produces a context score

$$Q_i \in [0, 1], \quad (3)$$

which quantifies the degree to which the case satisfies the contextual requirements encoded by the organization, institution, or application domain.

It is worth noting that the restriction of both predictive and contextual scores to the unit interval does not limit the generality of the proposed formulation. Any continuous score defined over an interval $[a, b]$ can be transformed into the normalized domain $[0, 1]$ through an affine mapping. This transformation preserves ordering relations and therefore maintains the ranking semantics required in prioritization tasks. Consequently, R_i and Q_i should be interpreted as normalized semantic scores rather than measurements expressed in a physical unit. The use of the unit interval allows the analysis of aggregation properties independently of the original scale of the predictive and contextual information. More importantly, the lower boundary $Q_i = 0$ represents a complete *contextual violation*.

The final prioritization score is then obtained through an aggregation operator

$$P_i = A(R_i, Q_i), \quad (4)$$

where A defines the integration semantics between predictive evidence and contextual assessment. Candidate cases are subsequently ranked according to P_i , and the corresponding Top- K subset is selected for intervention or resource allocation.

It is important to emphasize that the aggregation problem considered here differs from traditional fuzzy MCDM settings. In classical MCDM, aggregation operators combine multiple criteria describing different attributes of the same alternative, typically with the objective of capturing trade-offs among decision criteria. In CADEMAs-ML, however, the operator combines heterogeneous sources of decision evidence that play fundamentally different roles: a predictive assessment generated by cooperating machine-learning models and a contextual assessment derived from organizational, regulatory, or domain-specific knowledge.

Consequently, the aggregation operator does not merely determine the relative importance of criteria. Instead, it governs how predictive evidence interacts with contextual requirements and therefore acts as an explicit policy-design component within the decision architecture. The central question is not which operator produces the most desirable aggregate score in a multicriteria sense, but rather how alternative aggregation semantics affect policy compliance, contextual veto preservation, robustness to predictive overconfidence, and ranking stability under uncertainty.

Existing CADEMAs-ML applications employ a weighted linear combination of predictive and contextual scores [10]. However, the CADEMAs framework itself does not prescribe any particular aggregation operator [9]. Consequently, different aggregation semantics can be introduced without modifying either the predictive subsystem or the contextual subsystem. This modularity provides an ideal setting for studying aggregation as a policy-design mechanism and for analysing how different compensation regimes affect prioritization, policy compliance, robustness, and ranking stability. The following section introduces the aggregation operators considered in this study and develops the theoretical properties that motivate the subsequent experimental analysis.

4. Aggregation Semantics and Theoretical Analysis

4.1. Structural Properties of Aggregation Operators

Let

$$A : [0, 1]^2 \rightarrow [0, 1] \quad (5)$$

denote an aggregation operator that combines the predictive score R_i and the contextual score Q_i into the final prioritization score

$$P_i = A(R_i, Q_i). \quad (6)$$

Throughout this work, aggregation operators are interpreted as integration policies governing the interaction between predictive evidence and contextual knowledge within the CADEMAs-ML architecture. Rather than focusing on specific functional forms, we first introduce a set of structural properties that characterize different aggregation semantics.

Definition 1 (Boundedness). *An aggregation operator is bounded if*

$$0 \leq A(r, q) \leq 1, \quad \forall (r, q) \in [0, 1]^2. \quad (7)$$

Boundedness guarantees that the resulting prioritization score remains within the same normalized scale as the predictive and contextual assessments.

Definition 2 (Monotonicity). *An aggregation operator is monotonic if*

$$r_1 \leq r_2, q_1 \leq q_2 \implies A(r_1, q_1) \leq A(r_2, q_2). \quad (8)$$

Monotonicity ensures that improving either predictive evidence or contextual satisfaction cannot decrease the final prioritization score.

Definition 3 (Continuity). *An aggregation operator is continuous if small perturbations in its inputs produce arbitrarily small changes in the aggregated score.*

Continuity is particularly relevant when predictive and contextual scores are estimated from noisy data, since it prevents abrupt changes in prioritization due to infinitesimal variations of the inputs.

Definition 4 (Zero Absorption). *An aggregation operator satisfies the zero-absorption property if*

$$A(r, 0) = 0, \quad \forall r \in [0, 1]. \quad (9)$$

Under zero absorption, a null contextual assessment completely suppresses the final prioritization score regardless of the predictive evidence available.

Definition 5 (Contextual Veto Preservation). *An aggregation operator preserves contextual vetoes if*

$$Q_i = 0 \implies P_i = 0. \quad (10)$$

Within the CADEMAs-ML setting, contextual veto preservation means that contextual requirements cannot be overridden by predictive evidence. Observe that zero absorption is a sufficient condition for contextual veto preservation.

Definition 6 (Compensability). *An aggregation operator is compensatory if a reduction in one input can be offset by a sufficiently large increase in the other input while maintaining or improving the aggregated score.*

Compensability determines the extent to which strong predictive evidence may compensate for weak contextual support, or vice versa. At one extreme, fully compensatory operators permit unrestricted trade-offs between prediction and context. At the opposite extreme, non-compensatory operators prevent such trade-offs and therefore behave as hard contextual constraints.

These properties provide a theoretical basis for distinguishing alternative integration semantics within CADEMAs-ML. In particular, zero absorption and compensability play a central role in determining whether contextual requirements behave as soft preferences, partial constraints, or strict vetoes. The following subsection introduces the specific aggregation operators considered in this work and analyzes how they instantiate these structural properties.

4.2. Aggregation Operators Considered

The CADEMAs framework does not prescribe a specific mechanism for combining predictive and contextual assessments [9]. Consequently, different aggregation operators can be incorporated without modifying either the predictive subsystem or the contextual subsystem.

To investigate how aggregation semantics affect prioritization behavior, policy compliance, and robustness, this work considers three representative operators that span different levels of compensability between predictive evidence and contextual requirements.

1. *Linear aggregation.* The linear operator corresponds to the integration mechanism adopted in existing CADEMAs-ML implementations [10]:

$$A_L(R_i, Q_i) = \lambda R_i + (1 - \lambda)Q_i, \quad \lambda \in [0, 1]. \quad (11)$$

The parameter λ controls the relative importance assigned to predictive and contextual information. This operator is fully compensatory because reductions in one component can always be offset by sufficient increases in the other component.

2. *Weighted geometric aggregation.* As an intermediate alternative, we consider the weighted geometric mean:

$$A_G(R_i, Q_i) = R_i^\lambda Q_i^{1-\lambda}, \quad \lambda \in (0, 1). \quad (12)$$

Unlike the linear operator, the weighted geometric mean penalizes imbalanced evaluations and therefore reduces the degree of compensation between predictive and contextual information. Nevertheless, partial compensation remains possible whenever both inputs are strictly positive.

3. *Minimum aggregation.* Finally, we consider the minimum operator

$$A_M(R_i, Q_i) = \min(R_i, Q_i). \quad (13)$$

This operator represents a non-compensatory integration regime in which the final prioritization score cannot exceed the weakest component. Consequently, predictive evidence cannot compensate for insufficient contextual support.

Table 2 summarizes the structural properties of the three operators according to the definitions introduced in the previous subsection.

Table 2. Structural properties of the aggregation operators considered in this work.

Property	Linear	Geometric	Minimum
Boundedness	Yes	Yes	Yes
Monotonicity	Yes	Yes	Yes
Continuity	Yes	Yes	Yes
Zero absorption	No	Yes	Yes
Contextual veto preservation	No	Yes	Yes
Compensability	Full	Partial	No

These operators were selected because they represent three distinct integration semantics within CADEMAs-ML: unrestricted compensation, partial compensation, and strict contextual constraint. The following subsections analyze the theoretical consequences of these semantics and their implications for prioritization behavior.

4.3. Conditions for Contextual Veto Preservation

Within CADEMAs-ML, contextual information may represent different types of requirements, including organizational priorities, regulatory constraints, safety conditions, or expert-defined restrictions. In some applications, a contextual assessment should not behave as a soft preference but rather as a hard constraint. This situation occurs when a complete contextual violation must prevent an alternative from being prioritized, regardless of its predictive evidence.

This behavior can be characterized through the zero-absorption property introduced in Section 4.1. An aggregation operator provides contextual veto preservation if

$$A(r, 0) = 0, \quad \forall r \in [0, 1]. \quad (14)$$

Therefore, the question of whether CADEMAs-ML can enforce contextual vetoes reduces to determining whether the selected aggregation operator satisfies zero absorption.

Theorem 1 (Contextual veto preservation). *Let $A : [0, 1]^2 \rightarrow [0, 1]$ be an aggregation operator used in CADEMAs-ML. If A satisfies zero absorption, then every decision case i with a null contextual assessment,*

$$Q_i = 0, \quad (15)$$

receives a null prioritization score,

$$P_i = A(R_i, Q_i) = 0, \quad (16)$$

independently of its predictive score R_i .

Proof. The result follows directly from the definition of zero absorption. Since $Q_i = 0$, we have $P_i = A(R_i, Q_i) = A(R_i, 0)$. Because A satisfies zero absorption, $A(R_i, 0) = 0$, for every $R_i \in [0, 1]$. Therefore, $P_i = 0$. Hence, a complete contextual violation completely suppresses the final prioritization score regardless of the predictive evidence. $\square$

Corollary 1 (Veto preservation of the considered operators). *Among the aggregation operators considered in this work, the minimum and weighted geometric operators preserve contextual vetoes, whereas the linear operator does not.*

Proof. For the minimum operator, $A_M(r, 0) = \min(r, 0) = 0$. Similarly, for the weighted geometric operator, $A_G(r, 0) = r^\lambda 0^{1-\lambda} = 0$, for $\lambda \in (0, 1)$. In contrast, the linear operator yields $A_L(r, 0) = \lambda r$, which is greater than zero whenever $r > 0$ and $\lambda > 0$. Therefore, predictive evidence can override a complete contextual violation under linear aggregation. $\square$

The previous result establishes that contextual veto preservation is not a consequence of the predictive model or the contextual knowledge representation, but rather an emergent property of the integration semantics. Consequently, two CADEMAs-ML systems with identical predictive and contextual components may exhibit fundamentally different policy behavior depending only on the selected aggregation operator.

This observation motivates the subsequent analysis of rank reversal conditions and compensation thresholds, which characterize how much contextual improvement is required to offset predictive advantages under different aggregation regimes.

4.4. Conditions for Rank Reversal under Compensation

The previous section established that contextual veto preservation depends on the existence of a zero-absorption property. However, many aggregation operators used in decision-support systems follow a different philosophy: they allow partial compensation between the aggregated sources of evidence. In such cases, a sufficiently high value in one dimension may offset a lower value in another dimension.

This property is common in compensatory aggregation schemes, including weighted arithmetic means and several non-linear aggregation operators. Although compensation provides flexibility when predictive evidence and contextual information represent partially substitutable criteria, it also introduces the possibility that the final prioritization ranking differs from the ranking induced by predictive evidence alone.

To formalize this phenomenon, let $A : [0, 1]^2 \rightarrow [0, 1]$ be a continuous aggregation operator that is strictly increasing in both arguments. Given two alternatives i and j , assume

that alternative i has stronger predictive evidence, but alternative j has better contextual alignment: $R_i > R_j$ and $Q_i < Q_j$. The final prioritization scores are then $P_i = A(R_i, Q_i)$ and $P_j = A(R_j, Q_j)$. A rank reversal occurs when $P_j > P_i$, meaning that the contextual advantage of alternative j compensates for its predictive disadvantage.

Theorem 2 (General rank reversal condition under compensation). *Let $A : [0, 1]^2 \rightarrow [0, 1]$ be a continuous aggregation operator strictly increasing in both arguments and assume that it is compensatory, that is, there exist values $r_1 > r_2$ and $q_1 < q_2$ such that $A(r_2, q_2) > A(r_1, q_1)$.*

For two alternatives i and j satisfying $R_i > R_j$, a rank reversal in favour of alternative j occurs if and only if $Q_j > Q_i + \Phi_A(R_i, R_j, Q_i)$, where Φ_A is the minimum contextual improvement required to compensate for the predictive difference and is implicitly defined by $A(R_j, Q_i + \Phi_A) = A(R_i, Q_i)$.

Therefore, whenever the compensatory capacity of the operator is sufficient, a contextual advantage can overcome a predictive advantage.

Proof. Fix $R_i > R_j$ and Q_i . Define $g(q) = A(R_j, q) - A(R_i, Q_i)$. Because A is continuous and strictly increasing in its second argument, $g(q)$ is also continuous and strictly increasing. At $q = Q_i$, $g(Q_i) = A(R_j, Q_i) - A(R_i, Q_i) < 0$, because A is strictly increasing in its first argument and $R_j < R_i$. If the operator is compensatory, there exists a contextual value $q^* > Q_i$ such that $A(R_j, q^*) = A(R_i, Q_i)$. By continuity, this value is unique. Defining $\Phi_A(R_i, R_j, Q_i) = q^* - Q_i$, we obtain $P_j > P_i$ if and only if $Q_j > q^*$, which is equivalent to $Q_j > Q_i + \Phi_A(R_i, R_j, Q_i)$. $\square$

The theorem shows that rank reversal is not an artefact of a particular aggregation formula, but rather a consequence of allowing compensation between predictive evidence and contextual information. The specific threshold Φ_A depends on the semantics of each aggregation operator: operators with stronger compensation require smaller contextual improvements to reverse a predictive advantage, whereas operators with reduced compensation require larger contextual differences.

This distinction is particularly relevant in CADEMAs-ML because the two aggregated components do not represent equivalent evaluation criteria. The predictive score R_i summarizes evidence generated by machine learning models, whereas the contextual score Q_i represents compliance with organizational, regulatory, or domain-specific requirements. Therefore, the choice of aggregation operator implicitly defines the degree to which contextual requirements can modify or override data-driven recommendations.

4.5. Sensitivity to Predictive and Contextual Perturbations

5. Experimental Methodology

5.1. Objectives and Experimental Logic

5.2. Synthetic Population Generation

5.3. Monte Carlo Protocol

5.4. Evaluation Metrics

6. Simulation Results

6.1. Policy Compliance under Contextual Vetoes

6.2. Effects of Predictive Overconfidence

6.3. Robustness to Contextual Uncertainty

6.4. Sensitivity Analysis

7. Employee Attrition Case Study

7.1. Original CADEMAs-ML Application

7.2. Experimental Setup

7.3. Results and Interpretation

8. Discussion

8.1. Implications for Cooperative Automated Decision Making

8.2. Relation to Fuzzy MCDM and Aggregation Theory

8.3. Limitations

9. Conclusions and Future Work

Author Contributions: To be completed.

Funding: This research was supported by the project “Study, Analysis and Evaluation of Cooperative Automated Decision-Making Systems (CADEMAS)”, PID2023-146575NB-I00, funded by MCIU/AEI/I0.13039/501100011033 and by FSE+.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: To be completed.

Acknowledgments: Mariia Godz has a pre-doctoral research contract associated with the project “Study, Analysis and Evaluation of Cooperative Automated Decision-Making Systems (CADEMAS)”, PID2023-146575NB-I00, funded by MCIU/AEI/I0.13039/501100011033 and by FSE+.

Conflicts of Interest: The authors declare no conflicts of interest.

References

1. Teng, T.H.; Tan, A.H. Cognitive Agents Integrating Rules and Reinforcement Learning for Context-Aware Decision Support. In Proceedings of the Proceedings of the 2008 IEEE/WIC/ACM International Conference on Web Intelligence and Intelligent Agent Technology, 2008, pp. 318–321. <https://doi.org/10.1109/WIIAT.2008.163>.
2. Yao, W.; Kumar, A. Integrating Clinical Pathways into CDSS Using Context and Rules: A Case Study in Heart Disease. In Proceedings of the Proceedings of the 2nd ACM SIGHIT International Health Informatics Symposium, 2012, pp. 267–276. <https://doi.org/10.1145/2110363.2110431>.
3. Abu-Rasheed, H.; Abdulsalam, M.; Weber, C.; Fathi, M. Supporting Student Decisions on Learning Recommendations: An LLM-Based Chatbot with Knowledge Graph Contextualization for Conversational Explainability and Mentoring. *Frontiers in Artificial Intelligence* **2024**. <https://doi.org/10.35542/osf.io/ervym>.
4. Casado, D.F.; Fernandez, R.P. A Governance-Aware Methodological Framework for Cognitive AI Integration in Marine CAD: A Conceptual Pipe Routing Use Case. *Journal of Marine Science and Engineering* **2025**, *14*, 1352. <https://doi.org/10.3390/jmse14011352>.
5. Cronholm, S.; Göbel, H. Design Principles for Human-Centred AI. In Proceedings of the ECIS 2022 Research Papers, 2022. Research paper 32.

6. Fanfani, M.; Palesi, L.A.I.; Nesi, P. Human-Centered AI for Decision Support Systems: A Systematic Review of Application Domains, Architecture Designs, Current Trends and Future Directions. *Big Data and Cognitive Computing* **2025**, *10*, 186. <https://doi.org/10.3390/bdcc10060186>.
7. Smirnov, A.V.; Levashova, T.; Ponomarev, A.; Shilov, N. Methodology for Multi-Aspect Ontology Development: Ontology for Decision Support Based on Human-Machine Collective Intelligence. *IEEE Access* **2021**, *9*, 130721–130739. <https://doi.org/10.1109/ACCESS.2021.3116870>.
8. Marot, A.; Kelly, A.; Naglic, M.; Barbesant, V.; Cremer, J.L. Perspectives on Future Power System Control Centers for Energy Transition. *Journal of Modern Power Systems and Clean Energy* **2022**, *10*, 328–344. <https://doi.org/10.35833/mpce.2021.000673>.
9. Novoa-Hernández, P.; Pelta, D.; Godz, M.; Verdegay, J.; Buendia-Carrillo, D. CADEMÁS – A Framework for Cooperative Automated Decision-Making Systems. In Proceedings of the Proceedings of the 2026 IEEE Conference on Artificial Intelligence (CAI), 2026, pp. 756–761. <https://doi.org/10.1109/CAI68641.2026.11536392>.
10. Godz, M.; Novoa-Hernández, P.; Pelta, D. A Cooperative and Context-Aware Approach for Employee Attrition Prevention. In *Information Processing and Management of Uncertainty in Knowledge-Based Systems*; Springer Nature Switzerland, 2026; pp. 203–217. https://doi.org/10.1007/978-3-032-29000-7_15.
11. Yager, R. On Ordered Weighted Averaging Aggregation Operators in Multicriteria Decisionmaking. *IEEE Transactions on Systems, Man, and Cybernetics* **1988**, *18*, 183–190. <https://doi.org/10.1109/21.87068>.
12. Yager, R.R. Extending Multicriteria Decision Making by Mixing t-Norms and OWA Operators. *International Journal of Intelligent Systems* **2005**, *20*, 453–474. <https://doi.org/10.1002/int.20075>.
13. Mardani, A.; Nilashi, M.; Zavadskas, E.K.; Awang, S.; Zare, H.; Jamal, N.M. Decision making methods based on fuzzy aggregation operators: Three decades review from 1986 to 2017. *International Journal of Information Technology & Decision Making* **2018**, *17*, 391–466. <https://doi.org/10.1142/S021962201830001X>.
14. Csiszár, O. Ordered Weighted Averaging Operators: A Short Review. *IEEE Systems, Man, and Cybernetics Magazine* **2021**, *7*, 4–12. <https://doi.org/10.1109/MSMC.2020.3036378>.
15. del Moral, M.J.; Chiclana, F.; Tapia García, J.M.; Herrera-Viedma, E. An Analysis on Consensus Measures in Group Decision Making. In Proceedings of the Proceedings of the 2017 4th International Conference on Control, Decision and Information Technologies (CoDIT), 2017, pp. 0262–0267. <https://doi.org/10.1109/CoDIT.2017.8102605>.
16. Trillo, J.R.; Cabrerizo, F.J.; del Moral, M.J.; Morente-Molinera, J.A.; Tapia, J.M.; Herrera-Viedma, E. A Consensus-Based Multi-Criteria Group Decision-Making Method Based on an Aggregated Operator Customised by Experts. In Proceedings of the Proceedings of the 2023 9th International Conference on Control, Decision and Information Technologies (CoDIT), 2023. <https://doi.org/10.1109/CoDIT58514.2023.10284209>.
17. Yang, Y.W.; Lin, J.K.; Fu, Y.; Luo, Y. Tolerance Framework for Robust Group Multiple Criteria Decision Making. *Expert Systems with Applications* **2022**, *204*, 117430. <https://doi.org/10.1016/j.eswa.2022.118208>.
18. Liao, Z.; Liao, H.; Zhang, X. A Contextual Choquet Integral-Based Preference Learning Model Considering Both Criteria Interactions and the Compromise Effects of Decision-Makers. *Expert Systems with Applications* **2023**, *213*, 118977. <https://doi.org/10.1016/j.eswa.2022.118977>.
19. Stroppiana, D.; Boschetti, M.; Brivio, P.A.; Manfron, G.; Antoninetti, M. Explainable Multi-Criteria Data-Driven Environmental Status Assessment from Remote Sensing. In Proceedings of the Proceedings of the IEEE MELECON 2022, 2022. <https://doi.org/10.1109/MELECON53508.2022.9842919>.
20. Alcántud, J.C.R. Multi-Attribute Group Decision-Making Based on Intuitionistic Fuzzy Aggregation Operators Defined by Weighted Geometric Means. *Granular Computing* **2023**, *8*, 1857–1866. <https://doi.org/10.1007/s41066-023-00406-w>.
21. Merkert, J.; Mueller, M.; Hubl, M. A Survey of the Application of Machine Learning in Decision Support Systems. In Proceedings of the Proceedings of the 23rd European Conference on Information Systems (ECIS), 2015. <https://doi.org/10.18151/7217429>.
22. Pérez-Cañedo, B.; Porras, C.; Pelta, D.; Verdegay, J. Modeling Contexts as Fuzzy Propositions in Optimization Problems. *IEEE Transactions on Fuzzy Systems* **2023**, *31*, 1474–1483. <https://doi.org/10.1109/tfuzz.2022.3203786>.
23. Pérez-Cañedo, B.; Novoa-Hernández, P.; Porras, C.; Pelta, D.; Verdegay, J. Contextual analysis of solutions in a tourist trip design problem: A fuzzy logic-based approach. *Applied Soft Computing* **2024**, *154*, 111351. <https://doi.org/10.1016/j.asoc.2024.111351>.

Disclaimer/Publisher’s Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.