

# Beyond Linear Trade-offs in Context-Aware Decision Support: Non-Linear Aggregation and Rank Stability Analysis

Pavel Novoa-Hernández<sup>1,\*</sup> , Mariia Godz<sup>2</sup>  and David A. Pelta<sup>2</sup> 

<sup>1</sup> Dept. Computer Engineering and Systems, Universidad de La Laguna; pnovoah@ull.edu.es

<sup>2</sup> Dept. Computer Science and Artificial Intelligence, Universidad de Granada; {mariagodz,dpelta}@ugr.es

\* Correspondence: pnovoah@ull.edu.es

## Abstract

Hybrid automated decision-making systems increasingly combine machine-learning risk scores with contextual rules, expert knowledge, or organizational constraints. In such settings, the aggregation operator is not a technical detail: it determines whether predictive evidence may compensate for contextual non-compliance, or whether contextual requirements act as binding limits. Recent fuzzy multi-criteria decision-making research has produced many non-linear aggregation operators, but the literature provides little empirical evidence on their systemic effects, including policy violations, rank reversals, and the propagation of contextual noise. This paper studies these effects within the CADEMAs framework by comparing linear aggregation, the conjunctive minimum, and a weighted geometric operator. We first establish formal conditions under which aggregation semantics induce rank reversals or preserve predictive order under contextual homogeneity. We then run Monte Carlo simulations that stress-test resource allocation, targeted predictive overconfidence, and noisy contextual assessments. The results show that linear aggregation is efficient within a limited compliance region but can admit vetoed cases under high predictive weight, whereas the non-linear operators provide a structural firewall against zero-context alternatives. Under contextual noise, the minimum operator shows higher rank stability than the linear rule. Finally, a real employee-attrition case study from CADEMAs-ML illustrates how the same operators reorder borderline cases when both predictive risk and fuzzy context are strictly positive. The study supports operator selection as a governance decision rather than a purely mathematical modeling choice.

**Keywords:** automated decision-making; CADEMAs; fuzzy aggregation operators; multi-criteria decision making; rank stability; contextual governance; employee attrition

## 1. Introduction

Automated Decision-Making (ADM) systems are now used in settings where recommendations have economic, organizational, and ethical consequences. Machine learning has improved the ability to estimate risks and anticipate events, but a good prediction is not always an actionable decision. A hospital, a company, or a public agency may need to respect resource limits, institutional policies, safety constraints, or expert-defined priorities before acting on a model output. This is why many decision-support systems are moving from purely predictive pipelines toward hybrid architectures that combine data-driven evidence with rules, contextual information, and domain knowledge [1–3].

This shift creates a methodological problem that is easy to overlook. Once predictive evidence and contextual knowledge coexist, the final recommendation depends on how

Received:

Revised:

Accepted:

Published:

**Copyright:** © 2026 by the authors.

Submitted to *Mathematics* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](https://creativecommons.org/licenses/by/4.0/) license.

both signals are integrated. The aggregation rule therefore becomes part of the policy. It decides whether a high predicted risk can compensate for weak contextual support, whether a contextual veto must dominate the model, or whether the system should operate somewhere between those two extremes.

Novoa-Hernández et al. [4] proposed the Cooperative Automated Decision-Making System (CADEMAS) framework to make this separation explicit. CADEMAs organizes heterogeneous ADM units within a layered architecture: predictive or automated components produce partial evidence, an integration layer combines their outputs, and a contextual modulator evaluates the result under organizational or institutional criteria. In an employee-attribution instantiation, for example, machine-learning models estimate the probability that an employee will leave, while fuzzy contextual rules encode retention priorities such as strategic fit, development potential, or organizational vulnerability [5].

Existing CADEMAs instantiations, however, usually combine predictive risk and context through a linear convex rule. The appeal is clear: the parameter is interpretable, and the resulting score is easy to explain. Yet linearity also imposes a strong assumption. It treats prediction and context as fully compensatory dimensions. A sufficiently high risk score may offset poor contextual compatibility, and strong contextual support may offset weak predictive evidence. This is not necessarily wrong, but it is not neutral. In domains where contextual information represents a regulation, a safety requirement, or a fairness constraint, unrestricted compensation may lead to recommendations that are numerically plausible but institutionally unacceptable.

The fuzzy Multi-Criteria Decision-Making (MCDM) literature offers many alternatives to linear averaging. Ordered weighted, geometric, t-norm-based, Frank, Einstein, Aczel–Alsina, Dombi, and Diophantine fuzzy operators have been developed to represent different forms of interaction among criteria [6–14]. This literature is mathematically rich. It establishes properties such as boundedness, monotonicity, idempotency, and sensitivity to operator parameters. It also shows that aggregation choices can affect rankings in concrete decision problems.

What remains less developed is the systemic evaluation of these operators. Aires and Ferreira [15] show that rank reversal is a long-standing reliability problem in MCDM, and recent comparisons suggest that ranking stability can vary substantially across method architectures [16,17]. Still, the available evidence says little about the questions that matter for hybrid ADM systems: how often does an operator allow policy-violating cases into a resource-constrained top tier? Can a non-linear operator insulate a decision process from localized predictive overconfidence? How does fuzzy contextual noise propagate into final rankings? The literature has largely prioritized the construction of new operators and application examples over controlled benchmarking of such effects.

This paper addresses that gap for the CADEMAs setting. We do not propose another aggregation operator. Instead, we treat the aggregation stage as the mathematical place where governance semantics are encoded. A linear rule represents a fully compensatory policy; a conjunctive minimum represents a strict non-compensatory policy; a weighted geometric mean offers a non-linear intermediate regime with a hard zero boundary. We analyze how these regimes reshape the decision space, modify rankings, and respond to contextual uncertainty.

The empirical part of the paper has two layers. First, we run controlled Monte Carlo simulations to examine three aspects of operator behavior: policy compliance under increasing predictive weight, resilience to localized model overconfidence, and rank stability under noisy fuzzy context. Second, we apply the same operators to a real CADEMAs-ML employee-attribution cohort. The case study does not replace the simulations; it shows how

the same semantics appear in a concrete prioritization task where predictive risk and contextual alignment are both positive and heterogeneous.

The main contributions of this work are fivefold:

- We formulate the integration of predictive risk and fuzzy contextual knowledge within CADEMAs as a generalized aggregation problem.
- We study representative aggregation regimes with distinct compensation semantics: linear, conjunctive minimum, maximum, and weighted geometric integration.
- We characterize how aggregation semantics induce rank reversals or preserve predictive ordering under contextual homogeneity.
- We evaluate systemic effects that are underreported in the fuzzy MCDM literature: policy violations, targeted overconfidence, and noise propagation in prioritization rankings.
- We validate the interpretation on a real employee-attrition case study from CADEMAs-ML, showing how operator choice reorders borderline cases in a practical hybrid ADM workflow.

The remainder of the paper is organized as follows. Section 2 introduces the CADEMAs framework and its instantiations. Section 3 formalizes the aggregation models and their structural properties. Section 4 presents the simulation study and its findings. Section 5 applies the operators to a published employee-attrition instantiation. Section 6 summarizes the main conclusions, limitations, and directions for future work.

## 2. The CADEMAs Framework and Its Instantiations

Novoa-Hernández et al. [4] define CADEMAs (Cooperative Automated Decision-Making Systems) as a general framework for designing context-aware decision support systems in which heterogeneous Automated Decision-Making (ADM) components cooperate to produce a unified decision outcome. Unlike conventional ensemble methods, where multiple predictive models are combined solely to improve predictive performance, CADEMAs explicitly separates predictive inference from contextual governance. Predictive evidence and institutional knowledge are therefore treated as complementary sources of information that jointly determine the final recommendation.

Formally, a CADEMAs instance is represented by the tuple

$$S = \langle I, U, O, K, T \rangle, \quad (1)$$

where  $I$  denotes the Input Manager,  $U$  the set of Automated Decision-Making units,  $O$  the Decision Integration Layer,  $K$  the Contextual Modulator, and  $T$  the universe of admissible data representations. In practical terms,  $U$  is where machine-learning or rule-based models produce partial scores,  $O$  is where those scores are combined, and  $K$  is where organizational policy, expert knowledge, or fuzzy contextual rules enter the workflow. The architecture is intentionally generic and independent of any particular predictive model or contextual representation.

The ADM units generate partial decision signals from the available information, while the Decision Integration Layer harmonizes these heterogeneous outputs into a unified decision object. The Contextual Modulator then adjusts this preliminary outcome according to organizational policies, regulatory requirements, expert knowledge, or other contextual factors. This layered organization preserves a clear separation between what the data suggest and what the institution is willing to prioritize.

Recent applications have instantiated this architecture in domains where machine learning predictions must be reconciled with institutional decision policies. In the employee attrition scenario presented by Godz et al. [5], predictive models estimate the probability

that an employee voluntarily leaves the organization, whereas fuzzy contextual rules encode organizational priorities regarding employee retention. Similarly, in the SME insolvency risk application, predictive models estimate financial distress while contextual rules incorporate business and organizational considerations into the prioritization process.

Although these applications differ in domain and predictive models, they share a common decision structure. For each decision case  $i$ , the predictive subsystem produces a normalized risk score

$$R_i \in [0, 1], \quad (2)$$

which in CADEMAs represents the cooperative output of one or more ADM units (e.g., attrition probabilities aggregated across HR, finance, and operations models). The contextual module computes a normalized degree of contextual satisfaction

$$Q_i \in [0, 1], \quad (3)$$

which captures how strongly organizational policy supports prioritizing case  $i$  (e.g., strategic fit for retention under a digital-transformation context). Independently of the application domain, the integration layer must therefore construct a prioritization function

$$P_i = A(R_i, Q_i), \quad (4)$$

where  $A$  denotes an aggregation operator combining predictive evidence and contextual information. The final ranking over cases, and hence the Top- $K$  intervention list, follows directly from these scores.

Existing CADEMAs instantiations have adopted linear convex aggregation for  $A$ . In this work, we abstract the integration stage from any specific mechanism and study how different aggregation semantics reshape prioritization. The operator  $A$  becomes the central design choice through which CADEMAs encodes compensatory, veto-based, or intermediate governance policies.

### 3. Non-Linear Integration of Predictive Risk and Fuzzy Context

We now formalize how predictive risk and fuzzy contextual knowledge are combined within CADEMAs. We first examine the standard linear integration scheme, introduce representative non-linear alternatives, and characterize the ranking structure that determines which cases enter the prioritized tier.

#### 3.1. The Limits of Linear Convex Combination

Existing CADEMAs instantiations, including applications in employee attrition and SME insolvency risk assessment, typically integrate predictive and contextual information through a linear convex combination:

$$P_i = \lambda R_i + (1 - \lambda) Q_i, \quad \lambda \in [0, 1], \quad (5)$$

where  $\lambda$  controls the relative weight of the predictive layer in the final score. In CADEMAs-ML, this parameter is exposed to domain experts as a trade-off slider between model-driven risk and contextual alignment.

This formulation is fully compensatory. Consider two alternatives  $i$  and  $j$  such that  $R_i > R_j$  and  $Q_i < Q_j$ . Under linear aggregation, the ordering between  $i$  and  $j$  depends solely on  $\lambda$ . A deterioration in contextual compatibility can always be offset by an increase in predictive risk, and vice versa. For CADEMAs, this means that a case with weak contextual support may still rank above a better-aligned peer if its predicted risk is sufficiently high.

The compensatory structure follows from constant partial derivatives:

$$\frac{\partial P_i}{\partial R_i} = \lambda, \quad \frac{\partial P_i}{\partial Q_i} = 1 - \lambda,$$

so the marginal rate of substitution between prediction and context does not change with the scores themselves. Iso-priority contours in the  $(R_i, Q_i)$  plane are therefore straight lines. In settings where contextual information represents a safety requirement, a fairness constraint, or a hard organizational veto, unrestricted compensation may produce recommendations that are numerically high but institutionally unacceptable.

### 3.2. Non-Linear Aggregation in CADEMAs

To represent alternative governance logics, we consider a generalized aggregation operator

$$A : [0, 1]^2 \rightarrow [0, 1], \quad (6)$$

satisfying monotonicity and boundary conditions  $A(0, 0) = 0$  and  $A(1, 1) = 1$ . Each choice of  $A$  defines a different policy semantics for the CADEMAs integration layer. We distinguish three representative families.

1. *Conjunctive (non-compensatory) regime.* These operators model strict interaction between predictive and contextual dimensions, where poor performance in either dimension strongly penalizes the final score:

$$A_T^{\min}(R_i, Q_i) = \min(R_i, Q_i), \quad (7)$$

$$A_T^{\Pi}(R_i, Q_i) = R_i \cdot Q_i. \quad (8)$$

In CADEMAs, the minimum operator acts as a contextual gate: the final priority cannot exceed the weaker of the two signals. A case with  $Q_i = 0$  receives  $P_i = 0$  regardless of  $R_i$ , which corresponds to an institutional veto at the integration layer.

2. *Disjunctive (permissive) regime.* These operators capture settings in which satisfying either predictive or contextual criteria may be sufficient for prioritization:

$$A_S(R_i, Q_i) = \max(R_i, Q_i). \quad (9)$$

This regime emphasizes the most favorable signal and suits exploratory settings where either high risk or strong contextual alignment alone may justify attention.

3. *Non-linear compensatory regime.* Intermediate behavior between full compensation and strict non-compensation can be modeled through operators such as the weighted geometric mean:

$$A_G(R_i, Q_i) = R_i^\lambda Q_i^{1-\lambda}, \quad \lambda \in [0, 1]. \quad (10)$$

Unlike linear aggregation, this operator penalizes imbalance between  $R_i$  and  $Q_i$  while still allowing partial compensation. For CADEMAs, it offers a softer alternative to the minimum: contextual non-compliance reduces priority sharply but does not necessarily zero it unless  $Q_i = 0$ .

### 3.3. Rank Formulation and Structural Properties

Given a set of decision alternatives  $X = \{1, \dots, n\}$ , each alternative  $i \in X$  is assigned a score  $P_i = A(R_i, Q_i)$ , where  $R_i \in [0, 1]$  is the predictive risk and  $Q_i \in [0, 1]$  is the contextual satisfaction degree. We define the aggregated score vector  $\mathbf{P}(A) = (P_1, \dots, P_n)$  and the ranking function

$$\pi_A = \text{argsort}_{i \in X} P_i, \quad (11)$$

such that  $P_{\pi_A(1)} \geq P_{\pi_A(2)} \geq \dots \geq P_{\pi_A(n)}$ . The mapping  $A \rightarrow \mathbf{P}(A) \rightarrow \pi_A$  is the operational output of the CADEMAs integration layer: it determines which cases appear in the Top-K intervention tier and in what order.

This ranking  $\pi_A$  is the primary object of analysis in the subsequent sections. To interpret the empirical results, we establish two structural properties that govern how the choice of  $A$  affects the stability of  $\pi_A$  relative to the initial predictive ordering. We focus on *rank reversal*, the situation where the final ranking inverts the order suggested solely by predictive risk  $R$ .

Let  $i, j \in X$  be two alternatives such that the predictive layer identifies  $i$  as strictly riskier than  $j$  ( $R_i > R_j$ ). In a purely data-driven CADEMAs configuration,  $i$  would always rank above  $j$ . The following propositions characterize when contextual evaluation  $Q$  can reverse that order, depending on the semantic regime of  $A$ .

**Proposition 1** (Policy-Driven Veto and Forced Rank Reversal). *The following results characterize the conditions under which the aggregation mechanism induces rank reversals under linear and non-linear regimes:*

1. Under linear convex aggregation  $A_L$  with trade-off parameter  $\lambda \in (0, 1)$ , a rank reversal occurs ( $P_j > P_i$ ) if and only if the contextual advantage of alternative  $j$ , defined as  $\Delta Q = Q_j - Q_i$ , satisfies:

$$\Delta Q > \frac{\lambda}{1 - \lambda} (R_i - R_j) \quad (12)$$

2. Under the strict conjunctive aggregator  $A_T^{\min}$  (resp. disjunctive  $A_S^{\max}$ ), if the contextual score of the disadvantaged alternative is zero ( $Q_i = 0$ ) (resp. the advantaged alternative achieves perfect compliance,  $Q_j = 1$ ), a rank reversal is mathematically guaranteed, provided that the evaluation of the alternative not reaching the absorbing element is strictly positive (resp. strictly less than 1). Specifically, this requires  $R_j > 0$  and  $Q_j > 0$  for the conjunctive case, and  $\max(R_i, Q_i) < 1$  for the disjunctive case. The reversal holds independently of the predictive risk magnitude  $R_i$  (resp.  $R_j$ ).
3. Under the non-linear compensatory regime modeled by the weighted geometric mean  $A_G$ , assuming strictly positive inputs, a rank reversal occurs ( $P_j > P_i$ ) if and only if the logarithmic ratio of contextual advantage satisfies:

$$\ln\left(\frac{Q_j}{Q_i}\right) > \frac{\lambda}{1 - \lambda} \ln\left(\frac{R_i}{R_j}\right). \quad (13)$$

**Proof.** For part (1), we establish the equivalence by starting from the rank reversal condition:

$$P_j > P_i \iff \lambda R_j + (1 - \lambda)Q_j > \lambda R_i + (1 - \lambda)Q_i.$$

Rearranging the terms (a strictly reversible algebraic operation) yields:

$$\iff (1 - \lambda)(Q_j - Q_i) > \lambda(R_i - R_j).$$

Given  $\Delta Q = Q_j - Q_i$ , and since  $\lambda \in (0, 1)$  ensures  $(1 - \lambda) > 0$ , we can divide both sides by  $(1 - \lambda)$  without altering the inequality's direction or the logical equivalence:

$$\iff \Delta Q > \frac{\lambda}{1 - \lambda} (R_i - R_j).$$

This establishes both the necessary and sufficient conditions for the rank reversal under linear aggregation.



For part (2), we analyze the two dual regimes. In the conjunctive case, since  $Q_i = 0$ , we have  $P_i = \min(R_i, 0) = 0$ . Given  $R_j > 0$  and  $Q_j > 0$ , it follows that  $P_j = \min(R_j, Q_j) > 0$ , hence  $P_j > P_i$  regardless of the value of  $R_i$ . In the disjunctive case, since  $Q_j = 1$ , we have  $P_j = \max(R_j, 1) = 1$ . Given  $\max(R_i, Q_i) < 1$ , it follows that  $P_i = \max(R_i, Q_i) < 1$ , hence  $P_j > P_i$  regardless of the value of  $R_j$ .

Finally, for part (3), assuming  $R_i, R_j, Q_i, Q_j > 0$ , the rank reversal condition under the weighted geometric mean is:

$$P_j > P_i \iff R_j^\lambda Q_j^{1-\lambda} > R_i^\lambda Q_i^{1-\lambda}.$$

Taking the natural logarithm of both sides (a strictly monotonically increasing transformation that preserves the inequality) yields:

$$\iff \lambda \ln(R_j) + (1 - \lambda) \ln(Q_j) > \lambda \ln(R_i) + (1 - \lambda) \ln(Q_i).$$

Rearranging the terms gives:

$$\iff (1 - \lambda)(\ln(Q_j) - \ln(Q_i)) > \lambda(\ln(R_i) - \ln(R_j)).$$

Using logarithmic properties and dividing by  $(1 - \lambda) > 0$ , we obtain the desired equivalence:

$$\iff \ln\left(\frac{Q_j}{Q_i}\right) > \frac{\lambda}{1 - \lambda} \ln\left(\frac{R_i}{R_j}\right).$$

This completes the proof.  $\square$

The first part of Proposition 1 states the compensation boundary of linear CADEMAs integration: a low contextual score  $Q_i$  can always be offset by a sufficiently large predictive risk differential. Contextual signals then act as soft adjustments rather than binding constraints. The second part encodes the opposite logic. Under the minimum operator,  $Q_i = 0$  is an institutional veto that no amount of predictive evidence can overcome. The case is excluded from prioritization regardless of  $R_i$ , which protects policy compliance but may displace a high-risk alternative in the ranking.

The third part describes the geometric mean as an intermediate regime. The reversal threshold depends on score ratios rather than score differences, so compensation is asymmetric and scale-dependent. In CADEMAs terms, small contextual advantages can outweigh large risk differentials when both scores are low, whereas high scores require proportionally larger contextual gains to reverse the ranking. This offers a softer governance option between unrestricted linear trade-offs and hard veto semantics.

Together, these three regimes form a spectrum of integration policies that CADEMAs designers can encode through operator choice: fully compensatory ( $A_L$ ), strictly veto-based ( $A_T^{\min}$ ), or non-linear compensatory with a zero boundary ( $A_G$ ).

The second structural property addresses the complementary situation where contextual conditions are homogeneous across alternatives.

**Proposition 2** (Contextual Homogeneity and Predictive Rank Preservation). *Let  $Y \subseteq X$  be a subset of alternatives sharing a homogeneous contextual evaluation, such that  $Q_i = q$  for all  $i \in Y$ , with  $q > 0$ . For any aggregation operator  $A(R_i, Q_i)$  that is strictly monotonic with respect to its first argument when  $Q_i > 0$ , the strict ordinal ranking induced by the predictive model is fully preserved within  $Y$ . That is:*

$$P_i > P_j \iff R_i > R_j \quad \forall i, j \in Y. \quad (14)$$

**Proof.** Since  $Q_i = Q_j = q$  for any  $i, j \in Y$ , the aggregated scores are  $P_i = A(R_i, q)$  and  $P_j = A(R_j, q)$ . The direct implication ( $R_i > R_j \implies P_i > P_j$ ) follows directly from the strict monotonicity of  $A$  with respect to its first argument (with the second fixed at  $q > 0$ ). For the reverse implication ( $P_i > P_j \implies R_i > R_j$ ), we proceed by contraposition: if we assume  $R_i \leq R_j$ , the strict monotonicity condition would require  $A(R_i, q) \leq A(R_j, q)$ , meaning  $P_i \leq P_j$ . This contradicts our initial premise. Therefore, the double implication holds:

$$P_i > P_j \iff R_i > R_j$$

ensuring that the strict ordinal ranking induced by the predictive model is fully preserved.  $\square$

Proposition 2 guarantees that fuzzy contextual integration does not arbitrarily disrupt the predictive ranking when there is no contextual reason to differentiate between cases. If all alternatives in a group receive the same contextual score ( $Q_i = q$ ), the CADEMAs ranking within that group follows the ML risk hierarchy exactly. Contextual modulation then affects the absolute priority level but not the relative order among equally treated cases.

Together, these propositions establish the theoretical bounds for the empirical analysis in Section 4. Proposition 1 defines when context can override prediction under linear, non-linear compensatory, and veto-based regimes. Proposition 2 ensures that such overrides require genuine contextual differentiation rather than arising from the aggregation mechanism alone.

## 4. Simulation Study

Propositions 1 and 2 characterize when context can override prediction and when the predictive ranking should remain intact. These results are essential, but they do not yet tell us how aggregation semantics behave in a full CADEMAs prioritization workflow. In practice, the system must select a resource-constrained Top-K tier from a large population, cope with miscalibrated or overconfident predictive scores on specific groups, and integrate fuzzy contextual assessments that may be vague or noisy. We therefore examine three aspects of operator behavior:

1. **Policy compliance and predictive efficiency.** As the trade-off parameter  $\lambda$  increases, how far can each operator move toward a data-driven prioritization before institutionally vetoed cases ( $Q_i = 0$ ) enter the Top-K tier, and what predictive utility remains within the compliant region? On the basis of Proposition 1(1), we expect  $A_L$  to remain compliant only up to a finite  $\lambda$  threshold, whereas  $A_G$  should preserve zero violations across the full sweep because any zero context score forces  $P_i = 0$ .
2. **Resilience to localized overconfidence.** When predictive risk is artificially inflated on a vetoed sub-population, can the integration layer prevent false alarms from bypassing contextual policy? Proposition 1(2) suggests that  $A_T^{\min}$  and  $A_G$  should act as a structural firewall, while  $A_L$  should become vulnerable once predictive evidence dominates the final score.
3. **Rank stability under contextual noise.** How does imprecision in fuzzy contextual scores propagate into the final prioritization? We expect  $A_L$  to transmit noise directly into  $P_i$ , because every perturbation in  $Q_i$  enters the aggregated score linearly. The minimum operator may absorb part of this noise whenever predictive risk is the binding constraint, that is, when  $R_i \leq Q_i$ .

Each aspect is evaluated through Monte Carlo simulation under controlled populations designed to isolate one failure mode at a time. All curves report means over  $N = 30$  trials;



shaded bands indicate 95% trial-wise intervals. We describe differences as *structurally robust* when they hold in every trial (e.g.,  $\bar{V}_K = 0$  for non-linear operators under targeted overconfidence), and as *statistically supported* when mean trajectories diverge with non-overlapping interval bands.

#### 4.1. Experiment Settings

We simulate a universe of decision alternatives  $X = \{1, \dots, n\}$  with  $n = 1000$ . The system's goal is to select a prioritized subset of size  $K = \lfloor 0.10 \times n \rfloor = 100$ , representing a resource-constrained scenario where only the top 10% highest-priority alternatives receive intervention according to the ranking  $\pi_A$ . All simulations are implemented in Python 3; the source code is available in the project repository under `experiments/`.

Because each trial draws a fresh stochastic population (and, in the noise experiment, fresh perturbations), a single run could overfit to one random seed. We therefore repeat every experiment over  $N = 30$  independent Monte Carlo trials with seeds  $0, \dots, 29$ , holding  $n = 1000$  fixed within each trial. For any metric  $m$  (e.g.,  $V_K$ ,  $\bar{R}_K$ , or  $\tau$ ) evaluated at a given control setting, let  $m^{(1)}, \dots, m^{(N)}$  denote the trial-wise values. Figures plot the Monte Carlo mean  $\bar{m} = \frac{1}{N} \sum_{r=1}^N m^{(r)}$ ; shaded bands span the 2.5th and 97.5th percentiles across trials, forming a non-parametric 95% confidence interval. Qualitative claims about operator regimes (e.g., zero violations under  $Q_i = 0$  for  $A_T^{\min}$  and  $A_G$ ) refer to outcomes observed in all  $N$  trials; quantitative summaries for  $A_L$  are reported as  $(\bar{m}, \text{CI}_{95\%})$ .

#### Population designs.

The compliance and firewall experiments inject two sub-populations:

- **Standard sub-population ( $X_{\text{std}}$ , 80%):** Alternatives where  $R_i, Q_i \sim \mathcal{U}(0, 1)$ .
- **Veto sub-population ( $X_{\text{veto}}$ , 20%):** Alternatives with absolute policy violation ( $Q_i = 0$ ). In the compliance experiment, veto-group risk is elevated and fixed across the  $\lambda$  sweep ( $R_i \sim \mathcal{B}(8, 2)$ ). In the firewall experiment, the standard sub-population is held fixed across trials while veto-group risk is adversarially shifted as  $R_i \sim \mathcal{N}(\mu_{\text{shift}}, 0.05^2)$ , with  $\mu_{\text{shift}} \in [0.5, 1.0]$ .

For the noise experiment, we draw  $R_i, Q_i \sim \mathcal{U}(0, 1)$  for all  $n = 1000$  alternatives without injecting a veto sub-population. Gaussian perturbations are applied as  $Q'_i = \max(0, \min(1, Q_i + \epsilon))$  with  $\epsilon \sim \mathcal{N}(0, \sigma^2)$  and  $\sigma \in [0, 0.5]$ .

#### Operators and trade-off parameters.


We evaluate linear ( $A_L$ ), conjunctive minimum ( $A_T^{\min}$ ), and weighted geometric mean ( $A_G$ ), but not all three in every experiment:

- **Compliance:**  $A_L$  and  $A_G$  only;  $\lambda$  swept over  $[0.05, 0.95]$  (19 evenly spaced values).
- **Firewall:**  $A_L$ ,  $A_T^{\min}$ , and  $A_G$ ;  $\lambda = 0.85$  fixed. This data-driven setting is chosen because at  $\lambda = 0.5$  vetoed alternatives are capped at  $P_i \leq 0.5$ , masking the compensatory vulnerability of  $A_L$ .
- **Noise:**  $A_L$  and  $A_T^{\min}$  only;  $\lambda = 0.35$  fixed to grant the contextual layer sufficient weight so that noise in  $Q_i$  materially affects the ranking.

Multi-operator figures use a consistent visual encoding from the magma sequential palette ( $A_L$ ,  $A_T^{\min}$ ,  $A_G$  at colormap positions  $t \approx 0.28, 0.52, 0.76$ ), ordered linear  $\rightarrow$  min  $\rightarrow$  geometric in legends and panels.

#### Evaluation metrics.

Three operational metrics are computed trial-wise for the constrained Top-K scenario ( $K = 100$ ) and then aggregated across Monte Carlo replicates:



**Figure 1.** Opportunity cost of policy enforcement ( $N = 30$  Monte Carlo trials). Solid curves: mean; shaded bands: 95% trial-wise interval. (a) Policy violation rate:  $A_L$  remains robustly compliant ( $\bar{V}_K = 0$  in all trials) for  $\lambda \lesssim 0.75$ ; mean violations then rise sharply, reaching  $\bar{V}_K \approx 0.25$  at  $\lambda = 0.95$ .  $A_G$  maintains  $\bar{V}_K = 0$  across the full sweep. (b) Predictive efficiency: both operators achieve higher  $\bar{R}_K$  as  $\lambda$  increases within the compliant region; at  $\lambda = 0.80$ ,  $\bar{R}_K \approx 0.92$  for  $A_L$ ; at  $\lambda = 0.95$ ,  $A_G$  reaches  $\bar{R}_K \approx 0.93$  without violations.


- **Policy violation rate ( $V_K$ ):** The proportion of the prioritized Top-K tier belonging to the vetoed sub-population ( $X_{\text{veto}}$ ). In CADEMAs terms, this measures how often institutionally excluded cases still receive intervention priority. Reported as  $\bar{V}_K$  with 95% trial-wise intervals in Figures 1 and 2.
- **Predictive efficiency ( $\bar{R}_K$ ):** The average cooperative risk ( $R_i$ ) of the compliant standard cases selected in the Top-K tier. Higher values indicate better use of the predictive layer within the selected cohort (Figure 1b).
- **Rank stability under noise ( $\tau(\sigma)$ ):** Kendall's  $\tau$  between the ranking at  $\sigma = 0$  and the ranking induced after perturbing  $Q_i$  with noise level  $\sigma$ . Reported as  $\bar{\tau}(\sigma)$  with 95% trial-wise intervals in Figure 3.

#### 4.2. Policy Compliance and Predictive Efficiency

A robust CADEMAs instance must enforce institutional vetoes ( $V_K = 0$ ) without wasting intervention capacity on low-risk cases. Figure 1 traces  $\bar{V}_K$  and  $\bar{R}_K$  for  $A_L$  and  $A_G$  as  $\lambda$  varies over  $[0.05, 0.95]$ .

The results reveal a structural compliance ceiling in the linear regime. Contrary to the intuition that strict compliance requires down-weighting predictive evidence,  $A_L$  can keep  $\bar{V}_K \approx 0$  up to  $\lambda \approx 0.75$ – $0.80$ , and  $\bar{R}_K$  increases monotonically with  $\lambda$  over that range (from  $\approx 0.53$  at  $\lambda = 0.05$  to  $\approx 0.92$  at  $\lambda = 0.80$ ). The operational failure appears when the organization pushes further toward a data-driven regime: at  $\lambda = 0.95$ ,  $\bar{V}_K \approx 0.25$  (95% CI:  $[0.14, 0.34]$ ) even though  $\bar{R}_K \approx 0.95$ . This behavior matches the compensatory logic of Proposition 1(1): when  $\lambda$  is large, a sufficiently elevated  $R_i$  on a vetoed alternative can offset  $Q_i = 0$ , allowing policy-violating cases to enter the Top-K tier.

By contrast,  $A_G$  enforces a hard contextual boundary in CADEMAs: any  $Q_i = 0$  implies  $P_i = 0$  regardless of  $\lambda$ , so  $\bar{V}_K = 0$  in all 30 trials through  $\lambda = 0.95$ . The opportunity cost of linear integration is therefore not a permanent sacrifice of predictive utility at conservative  $\lambda$  values, but the inability to operate at high- $\lambda$ , high-utility settings without breaching organizational policy. This distinction is visible in the widening violation band



**Figure 2.** Algorithmic firewall under targeted ML overconfidence ( $\lambda = 0.85$ ;  $N = 30$  trials). Solid curves: mean; shaded bands: 95% trial-wise interval. Operators shown in order  $A_L$ ,  $A_T^{\min}$ ,  $A_G$ . For  $A_L$ ,  $\bar{V}_K$  stays near zero while  $\mu_{\text{shift}} \lesssim 0.75$ , then rises monotonically to  $\bar{V}_K \approx 0.30$  at  $\mu_{\text{shift}} = 1.0$  (95% CI:  $[0.10, 0.48]$ ).  $A_T^{\min}$  and  $A_G$  maintain  $\bar{V}_K = 0$  throughout; the bands collapse to the horizontal axis.

for  $A_L$  in Figure 1a while  $A_G$  remains flat at zero. Non-linear aggregation extends the compliant operating region available to hybrid ADM systems.


#### 4.3. Algorithmic Firewalling Under Targeted Overconfidence

Machine learning models frequently suffer from sub-population bias, generating localized false positive spikes. We inject an adversarial shift into  $X_{\text{veto}}$  and evaluate how effectively each aggregation operator prevents localized false alarms from bypassing the contextual policy.

As noted above,  $\lambda = 0.5$  would cap vetoed alternatives at  $P_i \leq 0.5$ , organically preventing them from entering the Top- $K$  subset ( $K = 100$ ) where the empirical cutoff typically exceeds 0.7. The fixed setting  $\lambda = 0.85$  therefore exposes the compensatory vulnerability of  $A_L$ . As shown in Figure 2, targeted ML overconfidence on the vetoed sub-population produces a monotonic escalation of  $\bar{V}_K$  for  $A_L$  once  $\mu_{\text{shift}} \gtrsim 0.80$ , whereas  $A_T^{\min}$  and  $A_G$  remain impervious in every trial.

Figure 2 shows the fragility of compensatory linearity under data-driven prioritization. As the ML layer becomes increasingly overconfident regarding the restricted cases ( $\mu_{\text{shift}} \rightarrow 1$ ),  $A_L$  allows inflated risk scores to override zero contextual compliance, with  $\bar{V}_K \approx 0.30$  on average. Nearly one third of the Top- $K$  tier then consists of institutionally vetoed alternatives. The wide shaded band for  $A_L$  reflects substantial trial-to-trial variability in how many vetoed cases breach the cutoff once  $\mu_{\text{shift}}$  exceeds  $\approx 0.80$ . In contrast, both non-linear regimes ( $A_T^{\min}$  and  $A_G$ ) maintain  $\bar{V}_K = 0$  throughout the adversarial shift in every trial.

This pattern aligns with Proposition 1: part (1) explains why  $A_L$  admits vetoed alternatives once predictive scores dominate ( $\lambda = 0.85$ ), whereas part (2) guarantees that  $A_T^{\min}$  assigns  $P_i = 0$  whenever  $Q_i = 0$ , and  $A_G$  enforces the same contextual boundary because  $Q_i = 0$  zeroes the aggregated score. For CADEMAs designers, these non-linear operators function as an algorithmic firewall that insulates contextual policy from localized predictive failures.



**Figure 3.** Rank stability under contextual noise ( $\lambda = 0.35$ ;  $N = 30$  trials). Solid curves: mean  $\bar{\tau}(\sigma)$ ; shaded bands: 95% trial-wise interval. At  $\sigma = 0.5$ ,  $\bar{\tau} \approx 0.40$  for  $A_L$  (95% CI:  $[0.38, 0.43]$ ) vs.  $\bar{\tau} \approx 0.47$  for  $A_T^{\min}$  (95% CI:  $[0.44, 0.50]$ ); the bands do not overlap, supporting a ranking-stability advantage for the conjunctive operator under moderate vagueness.

#### 4.4. Rank Stability Under Contextual Noise

Unlike crisp probabilistic outputs, contextual scores  $Q_i$  derived from fuzzy rules or human expert evaluations are inherently vague and subject to measurement noise. We measure rank stability as  $\bar{\tau}(\sigma)$ , the Monte Carlo mean of Kendall's  $\tau$  between the ranking at  $\sigma = 0$  and the ranking induced under each noise level.

Figure 3 reports  $\bar{\tau}(\sigma)$  for  $A_L$  and  $A_T^{\min}$  at  $\lambda = 0.35$ ; at  $\sigma = 0$ , both operators reproduce the baseline ranking exactly ( $\tau = 1.00$  in every trial).

The noise dynamics in Figure 3 extend the homogeneous baseline of Proposition 2. Once contextual homogeneity is broken by noise,  $A_L$  transmits perturbations globally: any fluctuation in  $Q_i$  directly perturbs  $P_i = \lambda R_i + (1 - \lambda)Q_i$ , producing a steep decline in  $\bar{\tau}$  (from 1.00 to  $\approx 0.40$  at  $\sigma = 0.5$ ). By contrast,  $A_T^{\min}$  exhibits selective insensitivity. Whenever  $R_i \leq Q_i$ , we have  $A_T^{\min}(R_i, Q_i) = R_i$ , so fluctuations in  $Q_i$  above the risk threshold leave  $P_i$  unchanged, yielding  $\bar{\tau} \approx 0.47$  at  $\sigma = 0.5$ . In CADEMAs terms, the minimum operator is more robust to vagueness in the contextual modulator when predictive risk already limits the final score.


## 5. A Real Example

The simulations above deliberately use controlled populations so that each failure mode can be observed in isolation. We now ask whether the same operator semantics remain visible in a concrete CADEMAs workflow. For this purpose, we reapply  $A_L$ ,  $A_T^{\min}$ , and  $A_G$  to the employee-attrition demo cohort ( $n = 20$ ) distributed with the CADEMAs-ML open-source tool<sup>1</sup>. The cohort is drawn from the IBM HR Analytics dataset and follows the attrition instantiation described by Godz et al. [5].

Predictive and contextual inputs.

Four H2O AutoML models (Culture and Wellbeing, Finance, Operations and Development, and Human Resources) produce case-level attrition probabilities that are AUC-weighted into a cooperative risk score  $R_i \in [0, 1]$ . Contextual alignment  $Q_i$  is computed

<sup>1</sup> [https://github.com/pnovo/cademas-app/tree/main/example\\_attrition](https://github.com/pnovo/cademas-app/tree/main/example_attrition).



**Figure 4.** Real attrition example ( $n = 20$ , Digital Transformation context,  $\lambda = 0.5$ ). **(a)** Cooperative risk  $R_i$  vs. context score  $Q_i$ ; highlighted cases illustrate contrasting risk–context profiles. **(b)** Top- $K$  prioritization scores ( $K = 2$ ) under  $A_L$ ,  $A_T^{\min}$ , and  $A_G$ .

from the *Digital / Technological Transformation* fuzzy specification (upskilling capacity, role adaptability, career scalability, and investment horizon), yielding  $Q_i \in [0.15, 0.71]$  on this cohort. The scores  $R_i$  and  $Q_i$  were exported once from CADEMAs-ML and vendored in the repository (experiments/data/attrition/); the attrition outcome label is retained for verification only and does not enter  $P_i$ .

#### Operator comparison.

We set  $\lambda = 0.5$  (balanced risk–context trade-off, as in Godz et al. [5]) and compare full-cohort rankings under the three integration regimes. Figure 4a locates each employee in the  $(Q_i, R_i)$  plane; Figure 4b reports the Top- $K$  prioritization tier with  $K = 2$  (10% of the cohort), mimicking a resource-constrained retention programme.

Table 1 reports the full prioritization landscape. At  $\lambda = 0.5$ , all three operators agree on the Top- $K$  tier: *Evelyn Taylor* ( $R_i = 0.96$ ,  $Q_i = 0.70$ ) and *Amelia Davis* ( $R_i = 0.99$ ,  $Q_i = 0.59$ ), both with confirmed attrition in the hold-out label. This consensus reflects jointly elevated risk and contextual alignment. It is the regime where Proposition 2 does not bind and operator choice is empirically secondary at the very top of the ranking.

Below the Top- $K$  cutoff, semantic differences emerge even without absolute policy vetoes ( $Q_i > 0$  for all cases). *Lucas Wright* carries the highest cooperative risk in the cohort ( $R_i = 0.99$ ) but ranks seventh under every operator because his contextual score is low ( $Q_i = 0.34$ ); under  $A_L$  and  $A_G$ , high  $R_i$  partially compensates ( $P_i = 0.66$  and  $0.58$ , respectively), whereas  $A_T^{\min}$  caps  $P_i$  at  $Q_i$ . The contrasting case is *Noah Lewis*, who attains the maximum contextual score ( $Q_i = 0.71$ ) yet minimal predictive risk ( $R_i = 0.07$ ). Linear integration yields  $P_i = 0.39$  (rank 11), while  $A_T^{\min}$  collapses to  $P_i = 0.07$  (rank 14) and  $A_G$  to  $0.22$  (rank 13): strong contextual alignment cannot lift a low-risk employee into the intervention tier when the integration layer is risk-aware or conjunctive. Intermediate rank reversals appear among high-risk attrition cases (e.g., *Benjamin Scott* moves from rank 5 under  $A_L$  to rank 3 under  $A_T^{\min}$  and  $A_G$  because his contextual score ( $Q_i = 0.46$ ) exceeds that of peers with similar  $R_i$ ).

These patterns complement the simulation results. The compliance and firewall experiments stress-test veto and overconfidence under synthetic  $Q_i = 0$  sub-populations, whereas this real cohort shows how compensation and conjunctive semantics reorder *borderline* cases when both  $R_i$  and  $Q_i$  are strictly positive. The rank-stability advantage of  $A_T^{\min}$  under contextual noise also has a qualitative analogue here: contextual fluctuations above the risk floor cannot inflate  $P_i$  when  $R_i \ll Q_i$ .

**Table 1.** Prioritization scores and ranks for the CADEMAs-ML attrition cohort ( $\lambda = 0.5$ , Digital Transformation context). Ranks are computed over all  $n = 20$  cases (1 = highest  $P_i$ ).

| Case             | $R_i$  | $Q_i$  | $P_i^{(L)}$ | $P_i^{(\min)}$ | $P_i^{(G)}$ | $r^{(L)}$ | $r^{(\min)}$ | $r^{(G)}$ | Attr. |
|------------------|--------|--------|-------------|----------------|-------------|-----------|--------------|-----------|-------|
| Evelyn Taylor    | 0.9555 | 0.7000 | 0.8278      | 0.7000         | 0.8178      | 1         | 1            | 1         | Yes   |
| Amelia Davis     | 0.9879 | 0.5914 | 0.7896      | 0.5914         | 0.7644      | 2         | 2            | 2         | Yes   |
| Theodore Adams   | 0.9803 | 0.4400 | 0.7101      | 0.4400         | 0.6568      | 3         | 5            | 5         | Yes   |
| William Young    | 0.9751 | 0.4450 | 0.7100      | 0.4450         | 0.6587      | 4         | 4            | 4         | Yes   |
| Benjamin Scott   | 0.9622 | 0.4550 | 0.7086      | 0.4550         | 0.6617      | 5         | 3            | 3         | Yes   |
| Charlotte Miller | 0.9204 | 0.4250 | 0.6727      | 0.4250         | 0.6254      | 6         | 6            | 6         | Yes   |
| Lucas Wright     | 0.9882 | 0.3350 | 0.6616      | 0.3350         | 0.5754      | 7         | 7            | 7         | Yes   |
| Harper Wilson    | 0.9641 | 0.2700 | 0.6170      | 0.2700         | 0.5102      | 8         | 8            | 8         | Yes   |
| Henry King       | 0.9676 | 0.1700 | 0.5688      | 0.1700         | 0.4056      | 9         | 10           | 9         | Yes   |
| Abigail Anderson | 0.9796 | 0.1500 | 0.5648      | 0.1500         | 0.3833      | 10        | 11           | 10        | Yes   |
| Noah Lewis       | 0.0651 | 0.7100 | 0.3876      | 0.0651         | 0.2150      | 11        | 14           | 13        | No    |
| Olivia Johnson   | 0.1234 | 0.6057 | 0.3646      | 0.1234         | 0.2734      | 12        | 12           | 11        | No    |
| Oliver Walker    | 0.0039 | 0.6886 | 0.3462      | 0.0039         | 0.0518      | 13        | 17           | 17        | No    |
| Emma Smith       | 0.0388 | 0.6420 | 0.3404      | 0.0388         | 0.1578      | 14        | 15           | 15        | No    |
| Sophia Brown     | 0.0037 | 0.5800 | 0.2918      | 0.0037         | 0.0463      | 15        | 18           | 18        | No    |
| Liam Clark       | 0.1713 | 0.3900 | 0.2806      | 0.1713         | 0.2585      | 16        | 9            | 12        | No    |
| Elijah Allen     | 0.0012 | 0.5600 | 0.2806      | 0.0012         | 0.0259      | 17        | 20           | 20        | No    |
| Mia Jones        | 0.0032 | 0.5200 | 0.2616      | 0.0032         | 0.0408      | 18        | 19           | 19        | No    |
| James Hall       | 0.0113 | 0.4200 | 0.2156      | 0.0113         | 0.0689      | 19        | 16           | 16        | No    |
| Ava Williams     | 0.1210 | 0.2275 | 0.1743      | 0.1210         | 0.1659      | 20        | 13           | 14        | No    |

6. Conclusions and Future Directions

This paper examined how the aggregation operator in CADEMAs encodes governance semantics when predictive risk and fuzzy contextual knowledge are combined. The theoretical analysis showed that linear integration is fully compensatory: a sufficiently large risk differential can override a weak contextual score (Proposition 1). The minimum and geometric operators instead enforce hard or softened contextual boundaries, and all monotonic operators preserve the predictive ranking when context is homogeneous across alternatives (Proposition 2).

The simulation study confirmed that these structural differences have operational consequences. Linear aggregation can remain policy-compliant up to moderate predictive weights while still improving predictive efficiency, but it admits vetoed cases into the Top-K tier once  $\lambda$  becomes large. Non-linear operators, especially  $A_T^{\min}$  and  $A_G$ , maintain zero violations under targeted overconfidence and extend the range of data-driven settings available without breaching institutional policy. Under contextual noise,  $A_T^{\min}$  preserves rankings more effectively than  $A_L$  when predictive risk already limits the final score. The employee-attrition case study showed that these semantics also matter below the Top-K cutoff in real workflows, where borderline cases are reordered according to compensation or conjunctive logic even when no absolute veto applies.

From a governance perspective, the choice of aggregation operator should be treated as a design decision, not a secondary modeling detail. A linear rule signals that predictive and contextual evidence may trade off freely. A minimum rule signals that contextual non-compliance is disqualifying. A geometric rule offers a non-linear compromise with a zero boundary. CADEMAs designers should align this choice with the institutional meaning of  $Q_i$  in each application domain.

Several limitations qualify these findings. The Monte Carlo experiments use synthetic populations with stylized veto and noise structures; they isolate failure modes but do not reproduce the joint distribution of real predictive and contextual scores. The attrition example uses a small cohort ( $n = 20$ ) and precomputed  $R_i$  and  $Q_i$  scores rather than retraining the H2O models within this repository. We compared three representative operators rather than the full family of fuzzy t-norms and OWA variants available in the MCDM literature. Finally, the analysis focuses on ranking and Top-K selection; it does not address downstream costs of false negatives, fairness constraints across demographic groups, or dynamic retraining of the predictive layer.



Future work should extend this line of inquiry in four directions. First, replicate the stress tests on additional CADEMAs instantiations and on larger real cohorts with documented policy vetoes. Second, broaden the operator benchmark to include other non-linear families and to study how parameter choice interacts with calibration of  $R_i$  and  $Q_i$ . Third, connect rank stability metrics to decision-theoretic costs (e.g., intervention budget, attrition losses, regulatory penalties) so that operator selection can be evaluated on outcome-relevant grounds. Fourth, integrate aggregation governance into CADEMAs-ML tooling, making operator semantics, compliance thresholds, and sensitivity to noise visible to domain experts during system configuration.

**Author Contributions:** For research articles with several authors, a short paragraph specifying their individual contributions must be provided. The following statements should be used “Conceptualization, X.X. and Y.Y.; methodology, X.X.; software, X.X.; validation, X.X., Y.Y. and Z.Z.; formal analysis, X.X.; investigation, X.X.; resources, X.X.; data curation, X.X.; writing—original draft preparation, X.X.; writing—review and editing, X.X.; visualization, X.X.; supervision, X.X.; project administration, X.X.; funding acquisition, Y.Y. All authors have read and agreed to the published version of the manuscript.”, please turn to the [CRediT taxonomy](#) for the term explanation. Authorship must be limited to those who have contributed substantially to the work reported.

**Funding:** Please add: “This research received no external funding” or “This research was funded by NAME OF FUNDER grant number XXX”. Check carefully that the details given are accurate and use the standard spelling of funding agency names at <https://search.crossref.org/funding>, any errors may affect your future funding.

**Institutional Review Board Statement:** Not applicable.

**Informed Consent Statement:** Not applicable.

**Data Availability Statement:** Simulation code and figure scripts are available in the project repository under `experiments/`. The real attrition example uses the CADEMAs-ML `example_attrition` bundle ([https://github.com/pnovoa/cademas-app/tree/main/example\\_attrition](https://github.com/pnovoa/cademas-app/tree/main/example_attrition)): case features (`cases_attrition.csv`), precomputed cooperative risk and fuzzy context scores (`experiments/data/attrit` and H2O MOJO models for score regeneration (`experiments/scripts/export_attrition_scores.py`). The underlying IBM HR Analytics dataset is publicly available on Kaggle.

**Acknowledgments:** In this section, you can acknowledge any support given that is not covered by the author contribution or funding sections. This may include administrative and technical support, or donations in kind (e.g., materials used for experiments). Where GenAI has been used for purposes such as generating text, data, or graphics, or for study design, data collection, analysis, or interpretation of data, please add “During the preparation of this manuscript/study, the author(s) used [tool name, version information] for the purposes of [description of use]. The authors have reviewed and edited the output and take full responsibility for the content of this publication.”

**Conflicts of Interest:** The authors declare no conflicts of interest.

## Appendix A

### Appendix A.1

The appendix is an optional section that can contain details and data supplemental to the main text—for example, explanations of experimental details that would disrupt the flow of the main text but nonetheless remain crucial to understanding and reproducing the research shown; figures of replicates for experiments of which representative data are shown in the main text can be added here if brief, or as Supplementary Materials. Mathematical proofs of results not central to the paper can be added as an appendix.

**Table A1.** This is a table caption.

| Title 1 | Title 2 | Title 3 |
|---------|---------|---------|
| Entry 1 | Data    | Data    |
| Entry 2 | Data    | Data    |

**Appendix B**

All appendix sections must be cited in the main text. In the appendices, figures, tables, etc., should be labeled starting with “A”, e.g., Figure A1, Figure A2, etc.

**References**

1. Kierner, S.; Kucharski, J.; Kierner, Z. Taxonomy of Hybrid Architectures Involving Rule-Based Reasoning and Machine Learning in Clinical Decision Systems: A Scoping Review. *Journal of Biomedical Informatics* **2023**, *144*, 104428. <https://doi.org/10.1016/j.jbi.2023.104428>.

2. Rahman, T.K.A.; Zainal, D.; Mubarak, H.; Shabir, F.; Anwar, N.; Asrowardi, I. Utilizing Machine Learning-Based Decision-Making to Align Higher Education Curriculum with Industry Requirements. *International Journal of Modern Education and Computer Science* **2025**, *17*, 1–15. <https://doi.org/10.5815/ijmecs.2025.04.01>.

3. Pérez-Domínguez, L.; Callejas-Cuervo, M.; García-Luna, F. Intuitionistic Fuzzy Dimensional Analysis in Multi-Criteria Decision-Making: A Computational Approach. *IEEE Access* **2024**, *12*, 1–15. <https://doi.org/10.1109/ACCESS.2024.3422616>.

4. Novoa-Hernández, P.; Pelta, D.A.; Godz, M.; Verdegay, J.L.; Buendia-Carrillo, D. CADEMÁS – A Framework for Cooperative Automated Decision-Making Systems. In Proceedings of the 2026 IEEE Conference on Artificial Intelligence (CAI), 2026, pp. 756–761. <https://doi.org/10.1109/CAI68641.2026.11536392>.

5. Godz, M.; Novoa-Hernández, P.; Pelta, D.A., A Cooperative and Context-Aware Approach for Employee Attrition Prevention. In *Information Processing and Management of Uncertainty in Knowledge-Based Systems*; Springer Nature Switzerland, 2026; pp. 203–217. [https://doi.org/10.1007/978-3-032-29000-7\\_15](https://doi.org/10.1007/978-3-032-29000-7_15).

6. Yager, R.R. On Ordered Weighted Averaging Aggregation Operators in Multicriteria Decisionmaking. *IEEE Transactions on Systems, Man, and Cybernetics* **1988**, *18*, 183–190. <https://doi.org/10.1109/21.87068>.

7. Grabisch, M.; Murofushi, T.; Sugeno, M. Aggregation Operators and Fuzzy Integrals. In *Fuzzy Measures and Integrals: Theory and Applications*; Physica-Verlag, 2000; pp. 3–41.

8. Tan, C.; Chen, X. Induced Intuitionistic Fuzzy Ordered Weighted Geometric Operator and Its Application to Multiple Attribute Decision Making. *International Journal of Intelligent Systems* **2010**, *25*, 1193–1209. <https://doi.org/10.1002/int.20446>.

9. Mahmood, T.; Ali, Z.; Baupradist, S.; Chinram, R. Complex Intuitionistic Fuzzy Aczel-Alsina Aggregation Operators and Their Application in Multi-Attribute Decision-Making. *Symmetry* **2022**, *14*, 2255. <https://doi.org/10.3390/sym14112255>.

10. Iampan, A.; García, G.S.; Riaz, M.; Athar Farid, H.M.; Chinram, R. Linear Diophantine Fuzzy Einstein Aggregation Operators for Multi-Criteria Decision-Making Problems. *Journal of Mathematics* **2021**, *2021*, 1–31. <https://doi.org/10.1155/2021/5548033>.

11. Riaz, M.; Farid, H.; Aslam, M.; Pamucar, D.; Božanić, D. Novel Approach for Third-Party Reverse Logistic Provider Selection Process under Linear Diophantine Fuzzy Prioritized Aggregation Operators. *Symmetry* **2021**, *13*, 1152. <https://doi.org/10.3390/sym13071152>.

12. Riaz, M.; Farid, H.; Wang, W.; Pamucar, D. Interval-Valued Linear Diophantine Fuzzy Frank Aggregation Operators with Multi-Criteria Decision-Making. *Mathematics* **2022**, *10*, 1811. <https://doi.org/10.3390/math10111811>.

13. Riaz, M.; Khokhar, M.A.; Pamucar, D.; Aslam, M. Cubic M-Polar Fuzzy Hybrid Aggregation Operators with Dombi’s T-Norm and T-Conorm with Application. *Symmetry* **2021**, *13*, 646. <https://doi.org/10.3390/sym13040646>.

14. Jeevitha, K.; Garg, H.; Vimala, J.; Aljuaid, H.; Abdel-Aty, A.H. Linear Diophantine Multi-Fuzzy Aggregation Operators and Its Application in Digital Transformation. *Journal of Intelligent & Fuzzy Systems* **2023**, *45*, 3097–3107. <https://doi.org/10.3233/JIFS-223844>.

15. Aires, R.F.d.F.; Ferreira, L. The Rank Reversal Problem in Multi-Criteria Decision Making: A Literature Review. *Pesquisa Operacional* **2018**, *38*, 331–362. <https://doi.org/10.1590/0101-7438.2018.038.02.0331>.

16. Amman, M.; Rashid, T.; Ali, A. Comparative Analysis of Multi-Criteria Decision-Making Methods Using q-Rung Orthopair Fuzzy Set: Evaluating the Impact of Distance Measures and Aggregation Operators. *Engineering Computations* **2025**, *42*, 3215–3239. <https://doi.org/10.1108/ec-04-2025-0375>.

17. Kizielewicz, B.; Bączkiewicz, A. A Hybrid MCDA Framework for Sustainability Assessment: Integrating Fuzzy Normalization and Weight Reidentification Based on External Rankings. *Procedia Computer Science* **2025**, *270*, 5665–5674. <https://doi.org/10.1016/j.procs.2025.10.035>.

**Disclaimer/Publisher’s Note:** The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.

609  
610  
611